

MCLENNAN COMMUNITY COLLEGE

Build Something

The MCC AI Sandbox · Tier 3

Ted Gonzalez, JD/MBA

You are not the beginner audience.

The gap is not capability. It is a sanctioned place to work and a boundary you can trust.

WHO IS IN THIS ROOM

What is actually missing.

HAVE

Fluency

You know what these tools are good at.
Most of you have a workflow you built yourself.

NEED

A place to work

Somewhere to experiment that is not production. Realistic data that is not real.

GET


Permission

Governance you can clear in ten minutes, a fictional college to break, and four sessions of build time.

Nationally, 56% of higher-ed employees already use non-institutional AI tools for work — EDUCAUSE, January 2026

THE CASE FOR THIS TIER

Campus tech leaders know where the value is.



Inside Higher Ed / Hanover 2026 Campus CTO-CIO Survey • n = 130, ±8 points

Almost nobody is doing this.

Not one named U.S. community college runs a sanctioned employee build-and-sandbox program. If MCC does, MCC is early.

THE OPPORTUNITY

Where the field actually is.

ZERO

Colleges doing it

No named U.S. community college with a staff build program.

SOME

Arguing for it

Times Higher Education, May 2026:
most institutional AI literacy programs
"stop at the same point: teaching staff
to use tools."

REAL

Companies already there

Checkr has HR, finance and legal staff
building their own workflow apps —
with tool stipends and engineers on call.

PART ONE

Governance first

Not paperwork at the end. The thing that makes the speed legitimate.

THE HONEST NUMBERS

0

of 5,600 production vite-coded applications scanned had basic security headers, CSRF protection, or properly scoped access.

Escape.tech scan, relayed by a Cloud Security Alliance research note, June 2026 — a security vendor scanning a sample of its own choosing. Treat the direction as sound, not the number.

AND WHEN IT GOES WRONG

Shadow AI has a price.

63%

No policy at all

Of organizations have no AI governance policy of any kind.

20%

Already breached

Had a shadow-AI-related breach by 2025.

\$670K

The premium

More per breach than at organizations with little or no shadow AI.

IBM, Cost of a Data Breach 2025

The sandbox is what lets you move fast.

The boundary is the permission — not the obstacle. Anyone who has worked in a regulated environment already knows this.

THE FRAME

What a boundary changes.

WITHOUT ONE

Every experiment is a judgment call made alone

At speed, by someone who does not want to be told no. So it never gets mentioned. That is how shadow IT is manufactured.

WITH ONE

You know in ten minutes whether you can proceed

Most of the time the answer is yes. The rest goes to IT early, when the idea is still cheap to change.

ONE PAGE. TEN MINUTES. REQUIRED BEFORE YOU BUILD.

The Tool Charter.

01 What does it do?

One sentence, no jargon. If you cannot, it is not scoped yet.

02 What data does it touch, and what class is it?

If anything is red — stop. IT and General Counsel review before you build.

03 Who else will use it?

Just me / my team / my division / students. Each step raises the bar.

04 Who owns it when you leave MCC?

A named person. If nobody signs for it, it does not get built.

05 What happens when it is wrong?

The failure mode, who notices, and the manual fallback.

TRIAGE

Four possible answers.

BUILD

Go

Green or yellow data, your own team, clear owner. This is most ideas — go, and stop asking.

CONDITIONS

Go, with limits

Fine to prototype, but not on real data, or not beyond your team. Named conditions, not vague caution.

REVIEW

IT and GC first

Touches red data, writes to a system of record, or needs a login. Bring it early.

DO NOT

Not this one

Some ideas should not be built by any of us. Hearing that is a successful outcome.

PART TWO

The sandbox

A fictional college, a data pack that looks real, and a simulator where nothing can break.

Sabine Crossing Community College does not exist.

Every roster, advising note, aid record, budget line and policy in the pack is fabricated.

Why invent a whole college?

WHY

The alternative is worse

Practising on toy examples teaches nothing. Practising on real records is the thing we are trying to prevent.

HOW

It looks real

Messy formats. Inconsistent program codes. An angry parent email. An aid file with an appeal in the student's own words.

SO

You can be reckless

Paste anything anywhere. Use a personal account. Break it. There is nothing behind it to protect.

What is in the pack.

01 **A roster that is a mess**

123 students, five spellings per program, GPAs above 4.0 it will average.

02 **A budget with planted errors**

Two planted errors. One divides by the wrong figure; one total stops short.

03 **A student handbook**

Nine sections — plus one listing what it deliberately does not cover.

04 **A scenario pack**

An angry parent email. Advising notes. An aid file. Eighteen sorting cards.

Practice the motions. Not the consequences.

A browser page that pretends to be a real working environment. Nothing connects to anything. Nothing can be damaged.

THE SIMULATOR

Six missions.

01

Ask, then doubt

02

Read before you click

03

Decline something

04

**The paste you almost
made**

05

Move around

06

Ship it

MISSION THREE

Decline something.

A connector called StudentSync Pro asks for full read and write on the student information system, every advising note including private ones, and all financial aid records — from a vendor with no agreement on file, with permission to train on the data and keep it after you disconnect.

Most people authorise it. That is not a character flaw — it is how these screens are designed.

The tell is not the permission list. It is the mismatch: a tool that produces "student success insights" has no business writing to financial aid.

MISSION FOUR

The paste you almost made.

You are asked to draft a financial aid appeal decision. The fastest possible path is to paste the student's record. The sandbox lets you reach for it, then stops you and explains exactly what would have gone wrong.

The impulse is real, and it arrives at 4:40 p.m. when you have one more letter to write.

The countermeasure is not willpower. It is a two-second pause and one question: does this task need the record, or just the policy? Nearly always, just the policy.

PART THREE

Now build something

Scoped assistants, spreadsheets, and a plain-English on-ramp to code.

Scoped assistants.

A reusable configuration with fixed instructions and a fixed set of source documents. It does one thing, consistently — and you do not need to write code to make one.

01

Instructions

Who it is for, what voice, and what to refuse.

02

Knowledge

The documents. Only the documents.
That is the trick.

03

Boundaries

Refusing well matters more than
answering well.

"If the answer is not in the provided documents, say so and direct the person to a named human."

Then test that it actually holds. Do not assume it does.

Break your neighbor's tool.

01 **It works on the questions its builder imagined**

You cannot test what you did not think of. Nobody can.

02 **So hand it over and tell them to break it**

Ask outside its sources. Ask something its role should refuse.

03 **Nobody's survives**

The gap between a prototype and a tool people can actually use.

04 **Fix the instruction, not the tool**

Most breakages are a missing sentence, not a missing feature.

WHERE MOST HOURS ACTUALLY LIVE

Spreadsheets, and the one rule.

RELIABLE

Explaining

A formula you inherited and nobody understands. Drafting a data dictionary. Cleaning inconsistent entries.

WITH CHECKING

Writing

A formula from a plain-English description. Restructuring a table. Suggesting a pivot approach.

DO NOT TRUST

Computing

Any number you have not spot-checked. The budget file has two planted errors and both look plausible.

Test every AI-generated formula on rows where you already know the answer.

THE ON-RAMP

Vibe coding: what is true, and what is not.

TRUE

Non-technical staff really do build

Checkr has HR, finance and legal people building workflow apps, supported by tool stipends and engineers on call.

ALSO TRUE

A third have replaced software

Of enterprises surveyed, 35% say they have replaced at least one SaaS product with something they built.

NOT TRUE

The manufacturing claim

That firms require one to two hours a day of vibe coding. No source. The only prominent figure runs the other way.

Rules for the build lab.

- 01 No red data. Not in the prototype, not "just for testing."**
Use the Sabine Crossing pack. It exists so this rule costs you nothing.
- 02 If it needs a login, it needs IT before anyone else uses it.**
Not before you build it. Before someone other than you touches it.
- 03 A prototype is not production. Say the words out loud.**
The most expensive mistake in this room is a demo that quietly becomes a dependency.
- 04 If you cannot explain it in one sentence, it is not scoped.**
Go back to question one of the charter.

THE PART EVERYONE SKIPS

Before you call it done.

01

Documented

What it does, what it assumes, what it does not handle. One page is plenty.

02

Owned

A named person who is not only you. If you left tomorrow, who answers for it?

03

Classified

Data class confirmed against the charter — including what it grew into while you built it.

04

Survivable

What breaks it, who notices, and what people do while it is down.

Programs like this have about a two-year half-life without sustainment. Design against that from day one.

AFTER THE LAST SESSION

Where this goes next.

WHO

AI Ambassadors

One per major unit, from this cohort, in a published directory. Be the person people ask before they paste something in.

WHEN

A monthly open lab

No agenda, all tools on the table, ninety minutes, same Friday every month. Almost no prep cost.

WHY

A showcase each semester

Demo what you built to anyone who shows up, leadership included. It doubles as the evidence file for continuation funding.

**You are not waiting
for the rules.
You are writing them.**

MCC has no institutional AI policy. What this room produces — the data chart, the tools list, the charter, the request form — is the first draft of the one it needs.